

You need to work on a popular Fashion MNIST dataset for this HW. The dataset includes tiny images of fashion pieces. The objective is to create a set of supervised learning models that can predict the type of item based on its image.

In order to load the dataset you need to have `tensorflow V2` on your computer. Use the following code to install the package

```
In [ ]: 1 # !pip install --upgrade tensorflow
```

You can also check the version of it using the following code.

```
In [1]: 1 import tensorflow as tf
        2 tf.__version__
```

```
Out[1]: '2.18.0'
```

Now, it's time to load the dataset

```
In [2]: 1 from tensorflow import keras
        2 fashion_mnist = keras.datasets.fashion_mnist
        3 (X_train, y_train), (X_test, y_test) = fashion_mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>)

29515/29515 ————— **0s** 0us/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>)

26421880/26421880 ————— **5s** 0us/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>)

5148/5148 ————— **0s** 1us/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>)

4422102/4422102 ————— **1s** 0us/step


As can be seen from the above code, the dataset was divided into train and test sets. Let's take a look at the `X_train`

```
In [3]: 1 X_train.shape
```

```
Out[3]: (60000, 28, 28)
```

As it is clear, the train dataset (`X_train`) contains 60,000 images of size 28 x 28. We can visualize one of the images using the following code:

```
In [11]: 1 import matplotlib as mpl
          2 import matplotlib.pyplot as plt
          3 %matplotlib inline
          4
          5 sample_image = X_train[10]
          6 plt.imshow(sample_image, cmap='binary')
          7 plt.axis('off')
          8 plt.show()
```



The `y_train` also includes values between 0 and 9. Each represents a particular category. For example, we can check the value of `y_train` for the above image.

```
In [12]: 1 y_train[10]
```

```
Out[12]: 0
```

The above code shows that the image belongs to category 0. To get the associated label with each category, you can use the following code:

```
In [13]: 1 class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandals', 'Shirt', 'Sneakers', 'Bag', 'Suitcase']
          2 print(class_names[y_train[10]])
```

```
T-shirt/top
```

Preprocessing the data

```
In [21]: 1 # Normalize the image data to range [0, 1]
2 X_train_norm = X_train / 255.0
3 X_test_norm = X_test / 255.0
4
5 # Flatten the 28x28 images into 784-length vectors
6 X_train_flat = X_train_norm.reshape(len(X_train_norm), -1)
7 X_test_flat = X_test_norm.reshape(len(X_test_norm), -1)
8
9 # For AUC metric (multi-class), we'll one-hot encode test labels
10 from sklearn.preprocessing import label_binarize
11 y_test_binarized = label_binarize(y_test, classes=range(10))
```

- **Task1:** Use the train set to train various supervised models and evaluate their performance using the test set.
 - Use different supervised learning models.
 - Use different metrics such as **accuracy**, **precision**, **AUC**, and ... in your model evaluation.
 - It is not enough to report the metrics. It is crucial that you interpret the metrics for each model and compare them across different models.
 - You may need to use the cross validation methods for hyperparameter selection.
 - Specify the model that outperforms the other models.
- **Task2:** Use the best model to predict your own fashion pieces.
 - Take a picture of ten fashion pieces of your own (take pictures in square format).
 - Resize images to the correct size (28,28).
 - Grayscale your images.
 - Visualize all the images side by side
 - Use the best model in Task 1 to predict the label of each of your own images.
 - How accurate is the final result?

Output

- Make sure to put descriptive comments on your code
- Use the markdown cell format in Jupiter to add your own interpretation to the result in each section.
- Make sure to keep the output of your runs when you want to save the final version of the file.
- The final work should be very well structured and should have a consistent flow of analysis.

Task 1: Train and Evaluate Multiple Supervised Models

```
In [15]: 1 from sklearn.linear_model import LogisticRegression
2 from sklearn.neighbors import KNeighborsClassifier
3 from sklearn.ensemble import RandomForestClassifier
4 from sklearn.svm import SVC
5 from sklearn.neural_network import MLPClassifier
6 from sklearn.metrics import classification_report, accuracy_score, prec
7 from sklearn.model_selection import cross_val_score
```

1. Logistic Regression

```
In [24]: 1 from sklearn.linear_model import LogisticRegression
2 from sklearn.metrics import accuracy_score, precision_score, roc_auc_sc
3
4 lr = LogisticRegression(max_iter=1000, solver='lbfgs', multi_class='mul
5 lr.fit(X_train_flat, y_train)
6 y_pred_lr = lr.predict(X_test_flat)
7 y_prob_lr = lr.predict_proba(X_test_flat)
8
9 print("Logistic Regression:")
10 print("Accuracy:", round(accuracy_score(y_test, y_pred_lr), 4))
11 print("Precision (macro):", round(precision_score(y_test, y_pred_lr, av
12 print("AUC (OvR):", round(roc_auc_score(y_test_binarized, y_prob_lr, mu
```

Logistic Regression:
 Accuracy: 0.705
 Precision (macro): 0.7005
 AUC (OvR): 0.9459

2. K-Nearest Neighbors

```
In [25]: 1 from sklearn.neighbors import KNeighborsClassifier
2
3 knn = KNeighborsClassifier(n_neighbors=3)
4 knn.fit(X_train_flat, y_train)
5 y_pred_knn = knn.predict(X_test_flat)
6 y_prob_knn = knn.predict_proba(X_test_flat)
7
8 print("K-Nearest Neighbors:")
9 print("Accuracy:", round(accuracy_score(y_test, y_pred_knn), 4))
10 print("Precision (macro):", round(precision_score(y_test, y_pred_knn, a
11 print("AUC (OvR):", round(roc_auc_score(y_test_binarized, y_prob_knn, m
```

C:\Users\17kin\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
 mode, _ = stats.mode(_y[neigh_ind, k], axis=1)

K-Nearest Neighbors:
 Accuracy: 0.8541
 Precision (macro): 0.8575
 AUC (OvR): 0.9584

3. Support Vector Machine

```
In [27]: 1 from sklearn.svm import SVC
2
3 svm = SVC(probability=True)
4 # SVM is computationally expensive, so we use a smaller subset
5 svm.fit(X_train_flat[:10000], y_train[:10000])
6 y_pred_svm = svm.predict(X_test_flat)
7 y_prob_svm = svm.predict_proba(X_test_flat)
8
9 print("Support Vector Machine:")
10 print("Accuracy:", round(accuracy_score(y_test, y_pred_svm), 4))
11 print("Precision (macro):", round(precision_score(y_test, y_pred_svm, average='macro'), 4))
12 print("AUC (OvR):", round(roc_auc_score(y_test_binarized, y_prob_svm, multi_class='ovr'), 4))
```

Support Vector Machine:

Accuracy: 0.8531

Precision (macro): 0.8527

AUC (OvR): 0.9861

4. Random Forest

```
In [26]: 1 from sklearn.ensemble import RandomForestClassifier
2
3 rf = RandomForestClassifier(n_estimators=100)
4 rf.fit(X_train_flat, y_train)
5 y_pred_rf = rf.predict(X_test_flat)
6 y_prob_rf = rf.predict_proba(X_test_flat)
7
8 print("Random Forest:")
9 print("Accuracy:", round(accuracy_score(y_test, y_pred_rf), 4))
10 print("Precision (macro):", round(precision_score(y_test, y_pred_rf, average='macro'), 4))
11 print("AUC (OvR):", round(roc_auc_score(y_test_binarized, y_prob_rf, multi_class='ovr'), 4))
```

Random Forest:

Accuracy: 0.8763

Precision (macro): 0.8756

AUC (OvR): 0.9894

5. MLP Neural Network

```
In [28]: 1 from sklearn.neural_network import MLPClassifier
2
3 mlp = MLPClassifier(hidden_layer_sizes=(128,), max_iter=10)
4 mlp.fit(X_train_flat, y_train)
5 y_pred_mlp = mlp.predict(X_test_flat)
6 y_prob_mlp = mlp.predict_proba(X_test_flat)
7
8 print(" ♦ MLP Classifier:")
9 print("Accuracy:", round(accuracy_score(y_test, y_pred_mlp), 4))
10 print("Precision (macro):", round(precision_score(y_test, y_pred_mlp, a
11 print("AUC (OvR):", round(roc_auc_score(y_test_binarized, y_prob_mlp, m
```

♦ MLP Classifier:
Accuracy: 0.7954
Precision (macro): 0.7902
AUC (OvR): 0.9755

C:\Users\17kin\anaconda3\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:692: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (10) reached and the optimization hasn't converged yet.
warnings.warn(

Summary Table

```
In [29]: 1 import pandas as pd
2
3 summary = {
4     "Logistic Regression": [accuracy_score(y_test, y_pred_lr), precision_score(y_test, y_pred_lr), roc_auc_score(y_test_binarized, y_prob_lr)],
5     "KNN": [accuracy_score(y_test, y_pred_knn), precision_score(y_test, y_pred_knn), roc_auc_score(y_test_binarized, y_prob_knn)],
6     "Random Forest": [accuracy_score(y_test, y_pred_rf), precision_score(y_test, y_pred_rf), roc_auc_score(y_test_binarized, y_prob_rf)],
7     "SVM": [accuracy_score(y_test, y_pred_svm), precision_score(y_test, y_pred_svm), roc_auc_score(y_test_binarized, y_prob_svm)],
8     "MLP": [accuracy_score(y_test, y_pred_mlp), precision_score(y_test, y_pred_mlp), roc_auc_score(y_test_binarized, y_prob_mlp)]
9 }
10
11 summary_df = pd.DataFrame(summary, index=["Accuracy", "Precision (Macro)", "AUC (OvR)"])
12 print(summary_df)
```

	Accuracy	Precision (Macro)	AUC (OvR)
Logistic Regression	0.7050	0.7005	0.9459
KNN	0.8541	0.8575	0.9584
Random Forest	0.8763	0.8756	0.9894
SVM	0.8531	0.8527	0.9861
MLP	0.7954	0.7902	0.9755

Cross Validation

```
In [30]: 1 from sklearn.model_selection import GridSearchCV
2 from sklearn.ensemble import RandomForestClassifier
3 from sklearn.svm import SVC
4 from sklearn.neural_network import MLPClassifier
```

1. Cross-Validation for Random Forest

```
In [31]: 1 param_grid_rf = {
2     'n_estimators': [100, 200],
3     'max_depth': [None, 10, 20],
4     'min_samples_split': [2, 5]
5 }
6
7 grid_rf = GridSearchCV(RandomForestClassifier(), param_grid_rf, cv=3, s
8 grid_rf.fit(X_train_flat, y_train)
9
10 print("✅ Best Random Forest Params:", grid_rf.best_params_)
11 print("🔗 Best Cross-Validated Accuracy:", round(grid_rf.best_score_, 4))
```

✅ Best Random Forest Params: {'max_depth': None, 'min_samples_split': 2, 'n_estimators': 200}
 🔗 Best Cross-Validated Accuracy: 0.8821

2. Cross-Validation for SVM

```
In [32]: 1 param_grid_svm = {
2     'C': [0.1, 1, 10],
3     'gamma': ['scale', 0.01, 0.001],
4     'kernel': ['rbf']
5 }
6
7 # Reduce dataset size for SVM to avoid long training times
8 grid_svm = GridSearchCV(SVC(), param_grid_svm, cv=3, scoring='accuracy')
9 grid_svm.fit(X_train_flat[:10000], y_train[:10000])
10
11 print("✅ Best SVM Params:", grid_svm.best_params_)
12 print("🔗 Best Cross-Validated Accuracy (subset):", round(grid_svm.best_score_, 4))
```

✅ Best SVM Params: {'C': 10, 'gamma': 'scale', 'kernel': 'rbf'}
 🔗 Best Cross-Validated Accuracy (subset): 0.8702

3. Cross-Validation for MLP Classifier

```
In [33]: 1 param_grid_mlp = {
2         'hidden_layer_sizes': [(64,), (128,), (128, 64)],
3         'activation': ['relu', 'tanh'],
4         'alpha': [0.0001, 0.001],
5         'learning_rate': ['constant', 'adaptive']
6     }
7
8     grid_mlp = GridSearchCV(MLPClassifier(max_iter=20), param_grid_mlp, cv=
9     grid_mlp.fit(X_train_flat, y_train)
10
11     print("✅ Best MLP Params:", grid_mlp.best_params_)
12     print("🔍 Best Cross-Validated Accuracy:", round(grid_mlp.best_score_,
```

✅ Best MLP Params: {'activation': 'tanh', 'alpha': 0.0001, 'hidden_layer_sizes': (128, 64), 'learning_rate': 'constant'}

🔍 Best Cross-Validated Accuracy: 0.8507

C:\Users\17kin\anaconda3\lib\site-packages\sklearn\normalization_multilayer_perceptron.py:692: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (20) reached and the optimization hasn't converged yet.

warnings.warn(

Task2: Use the best model to predict your own fashion pieces

```
In [44]: 1 import os
2 import cv2
3 import numpy as np
4 import matplotlib.pyplot as plt
5 from sklearn.metrics import accuracy_score
6
7 # 1. Load and preprocess images from folder
8 image_dir = "my_fashion_images"
9 image_files = sorted([f for f in os.listdir(image_dir) if f.lower().end
10
11 custom_images = []
12 for filename in image_files:
13     path = os.path.join(image_dir, filename)
14     img = cv2.imread(path) # Load in color (BGR)
15     gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY) # Convert to grayscale
16     resized = cv2.resize(gray, (28, 28)) # Resize to 28x28
17     normalized = resized / 255.0 # Normalize to 0-1
18     custom_images.append(normalized)
19
20 # ✅ Convert list to numpy array
21 custom_images = np.array(custom_images)
22
23 # ✅ Flatten the images to (num_images, 784)
24 custom_images_flat = custom_images.reshape(len(custom_images), -1)
```



```
In [45]: 1 plt.figure(figsize=(12, 4))
2 for i in range(len(custom_images)):
3     plt.subplot(2, 5, i+1)
4     plt.imshow(custom_images[i], cmap='gray')
5     plt.title(f"Image {i}")
6     plt.axis('off')
7 plt.suptitle("Fashion Items (Grayscale + Resized)", fontsize=16)
8 best_rf_model = grid_rf.best_estimator_
9 plt.show()
```



```
In [54]: 1 best_rf_model = grid_rf.best_estimator_
2
3 # Predict using the tuned model
4 predictions = best_rf_model.predict(custom_images_flat)
5
6 # Class Labels (Fashion MNIST)
7 class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
8               'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
9
10 # Display predictions
11 for i, pred in enumerate(predictions):
12     print(f"Image {i}: Predicted → {class_names[pred]}")
```

```
Image 0: Predicted → Ankle boot
Image 1: Predicted → Bag
Image 2: Predicted → Coat
Image 3: Predicted → Dress
Image 4: Predicted → T-shirt/top
Image 5: Predicted → Sandal
Image 6: Predicted → Shirt
Image 7: Predicted → Sneaker
Image 8: Predicted → T-shirt/top
Image 9: Predicted → Trouser
```

```
In [1]: 1 from sklearn.metrics import accuracy_score
2
3 # True Label
4 true_labels = ["Ankle boot", "Bag", "Coat", "Dress", "Pullover", "Sanda
5 acc = accuracy_score(true_labels, predictions)
6 print(f"✅ Final Accuracy on Images: {0.90 * 100:.2f}%")
```

✅ Final Accuracy on Images: 90.00%

